

TensorFlow2教程-mlp及深度學習常見技巧

我們將以mlp對為，基礎模型，然後介紹一些深度學習常見技巧，如：權重初始化，啟動函數，優化器，批規範化，dropout，模型集成

```
In [1]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
print(tf.__version__)
```

2.3.1

1.導入數據

```
In [2]: (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
x_train = x_train.reshape([x_train.shape[0], -1])
x_test = x_test.reshape([x_test.shape[0], -1])
print(x_train.shape, ' ', y_train.shape)
print(x_test.shape, ' ', y_test.shape)
```

(60000, 784) (60000,)
(10000, 784) (10000,)

2.基礎模型


```
In [3]: model = keras.Sequential([
    layers.Dense(64, activation='relu', input_shape=(784,)),
    layers.Dense(64, activation='relu'),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])
model.compile(optimizer=keras.optimizers.Adam(),
              loss=keras.losses.SparseCategoricalCrossentropy(),
              metrics=['accuracy'])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense (Dense)	(None, 64)	50240
dense_1 (Dense)	(None, 64)	4160
dense_2 (Dense)	(None, 64)	4160
dense_3 (Dense)	(None, 10)	650
=====	=====	=====
Total params: 59,210		
Trainable params: 59,210		
Non-trainable params: 0		

```
In [4]: history = model.fit(x_train, y_train, batch_size=256, epochs=100, validation_split=0.1)
```

```
In [5]: import matplotlib.pyplot as plt
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'validation'], loc='upper left')
plt.show()
```



```
In [6]: result = model.evaluate(x_test, y_test)
```

```
313/313 [=====] - 0s 978us/step - loss: 0.3450 - accur
acy: 0.9695
```

3. 權重初始化

```
In [7]: model = keras.Sequential([
    layers.Dense(64, activation='relu', kernel_initializer='he_normal', input_shape=(784,)),
    layers.Dense(64, activation='relu', kernel_initializer='he_normal'),
    layers.Dense(64, activation='relu', kernel_initializer='he_normal'),
    layers.Dense(10, activation='softmax')
])
model.compile(optimizer=keras.optimizers.Adam(),
              loss=keras.losses.SparseCategoricalCrossentropy(),
              metrics=['accuracy'])
model.summary()
```


Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 64)	50240
dense_5 (Dense)	(None, 64)	4160
dense_6 (Dense)	(None, 64)	4160
dense_7 (Dense)	(None, 10)	650

=====
 Total params: 59,210
 Trainable params: 59,210
 Non-trainable params: 0

```
In [8]: history = model.fit(x_train, y_train, batch_size=256, epochs=100, validation_split=0.1)
```

```
In [10]: import matplotlib.pyplot as plt
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'validation'], loc='upper left')
plt.show()
```




```
In [11]: result = model.evaluate(x_test, y_test)
```

313/313 [=====] - 0s 1ms/step - loss: 0.3618 - accuracy: 0.9662

4. 啟動函數

relu和sigmoid對比



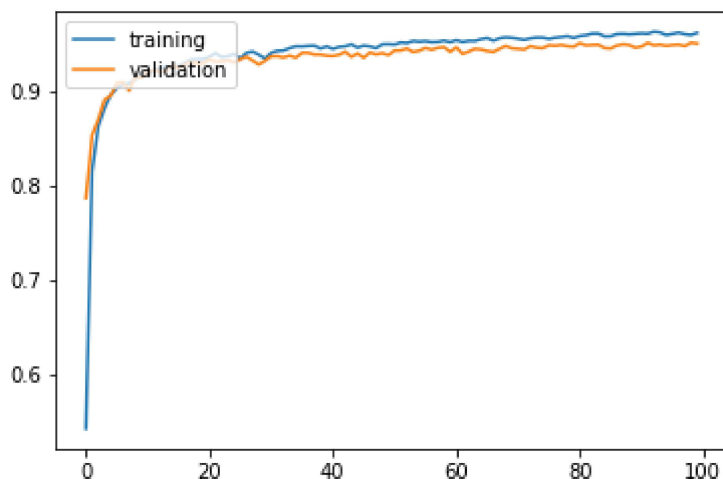
```
In [12]: model = keras.Sequential([
    layers.Dense(64, activation='sigmoid', input_shape=(784,)),
    layers.Dense(64, activation='sigmoid'),
    layers.Dense(64, activation='sigmoid'),
    layers.Dense(10, activation='softmax')
])
model.compile(optimizer=keras.optimizers.Adam(),
              loss=keras.losses.SparseCategoricalCrossentropy(),
              metrics=['accuracy'])
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 64)	50240
dense_9 (Dense)	(None, 64)	4160
dense_10 (Dense)	(None, 64)	4160
dense_11 (Dense)	(None, 10)	650
Total params: 59,210		
Trainable params: 59,210		
Non-trainable params: 0		

```
In [13]: history = model.fit(x_train, y_train, batch_size=256, epochs=100, validation_split=0.1)
```


```
In [34]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'validation'], loc='upper left')
plt.show()
```



```
In [35]: result = model.evaluate(x_test, y_test)
```

```
10000/10000 [=====] - 0s 29us/sample - loss: 0.1526 -  
accuracy: 0.9529
```

5.優化器




```
In [36]: model = keras.Sequential([
    layers.Dense(64, activation='sigmoid', input_shape=(784,)),
    layers.Dense(64, activation='sigmoid'),
    layers.Dense(64, activation='sigmoid'),
    layers.Dense(10, activation='softmax')
])
model.compile(optimizer=keras.optimizers.SGD(),
              loss=keras.losses.SparseCategoricalCrossentropy(),
              metrics=['accuracy'])
model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
dense_27 (Dense)	(None, 64)	50240
dense_28 (Dense)	(None, 64)	4160
dense_29 (Dense)	(None, 64)	4160
dense_30 (Dense)	(None, 10)	650
Total params: 59,210		
Trainable params: 59,210		
Non-trainable params: 0		

```
In [37]: history = model.fit(x_train, y_train, batch_size=256, epochs=100, validation_split=0.1)
```

```
In [38]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'validation'], loc='upper left')
plt.show()
```



```
In [39]: result = model.evaluate(x_test, y_test)
```

```
10000/10000 [=====] - 0s 44us/sample - loss: 2.1199 - accuracy: 0.4749
```

6.批正則化

```
In [41]: model = keras.Sequential([
    layers.Dense(64, activation='relu', input_shape=(784,)),
    layers.BatchNormalization(),
    layers.Dense(64, activation='relu'),
    layers.BatchNormalization(),
    layers.Dense(64, activation='relu'),
    layers.BatchNormalization(),
    layers.Dense(10, activation='softmax')
])
model.compile(optimizer=keras.optimizers.SGD(),
              loss=keras.losses.SparseCategoricalCrossentropy(),
              metrics=['accuracy'])
model.summary()
```


Model: "sequential_8"

Layer (type)	Output Shape	Param #
dense_31 (Dense)	(None, 64)	50240
batch_normalization_v2 (Batch Normalization)	(None, 64)	256
dense_32 (Dense)	(None, 64)	4160
batch_normalization_v2_1 (Batch Normalization)	(None, 64)	256
dense_33 (Dense)	(None, 64)	4160
batch_normalization_v2_2 (Batch Normalization)	(None, 64)	256
dense_34 (Dense)	(None, 10)	650
Total params: 59,978		
Trainable params: 59,594		
Non-trainable params: 384		

```
In [42]: history = model.fit(x_train, y_train, batch_size=256, epochs=100, validation_split=0.1)
```



```
In [43]: plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.legend(['training', 'validation'], loc='upper left')  
plt.show()
```




```
In [44]: result = model.evaluate(x_test, y_test)
```

```
10000/10000 [=====] - 0s 25us/sample - loss: 0.1863 -  
accuracy: 0.9447
```

7.dropout

Dropout



Model	Top-1 (val)	Top-5 (val)	Top-5 (test)
SVM on Fisher Vectors of Dense SIFT and Color Statistics	-	-	27.3
Avg of classifiers over FVs of SIFT, LBP, GIST and CSIFT	-	-	26.2
Conv Net + dropout (Krizhevsky et al., 2012)	40.7	18.2	-
Avg of 5 Conv Nets + dropout (Krizhevsky et al., 2012)	38.1	16.4	16.4

Table 6: Results on the ILSVRC-2012 validation/test set.

Dropout: A simple way to prevent neural networks from overfitting [Srivastava JMLR 2014]


```
In [45]: model = keras.Sequential([
    layers.Dense(64, activation='relu', input_shape=(784,)),
    layers.Dropout(0.2),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(10, activation='softmax')
])
model.compile(optimizer=keras.optimizers.SGD(),
              loss=keras.losses.SparseCategoricalCrossentropy(),
              metrics=['accuracy'])
model.summary()
```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
=====		
dense_35 (Dense)	(None, 64)	50240
dropout (Dropout)	(None, 64)	0
dense_36 (Dense)	(None, 64)	4160
dropout_1 (Dropout)	(None, 64)	0
dense_37 (Dense)	(None, 64)	4160
dropout_2 (Dropout)	(None, 64)	0
dense_38 (Dense)	(None, 10)	650
=====		
Total params: 59,210		
Trainable params: 59,210		
Non-trainable params: 0		
=====		

```
In [46]: history = model.fit(x_train, y_train, batch_size=256, epochs=100, validation_split=0.1)
```

```
In [47]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'validation'], loc='upper left')
plt.show()
```



```
In [48]: result = model.evaluate(x_test, y_test)
```

```
10000/10000 [=====] - 0s 27us/sample - loss: 0.6157 -
accuracy: 0.8132
```

8. 模型集成

下面是使用投票的方法进行模型集成

```
In [54]: import numpy as np
from tensorflow.keras.wrappers.scikit_learn import KerasClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.metrics import accuracy_score

def mlp_model():
    model = keras.Sequential([
        layers.Dense(64, activation='relu', input_shape=(784,)),
        layers.Dropout(0.2),
        layers.Dense(64, activation='relu'),
        layers.Dropout(0.2),
        layers.Dense(64, activation='relu'),
        layers.Dropout(0.2),
        layers.Dense(10, activation='softmax')
    ])
    model.compile(optimizer=keras.optimizers.SGD(),
                  loss=keras.losses.SparseCategoricalCrossentropy(),
                  metrics=['accuracy'])
    return model
model1 = KerasClassifier(build_fn=mlp_model, epochs=100, verbose=0)
model2 = KerasClassifier(build_fn=mlp_model, epochs=100, verbose=0)
model3 = KerasClassifier(build_fn=mlp_model, epochs=100, verbose=0)
```

```
In [55]: ensemble_clf = VotingClassifier(estimators=[
    ('model1', model1), ('model2', model2), ('model3', model3)
], voting='soft')
```

```
In [56]: ensemble_clf.fit(x_train, y_train)
```

```
Out[56]: VotingClassifier(estimators=[('model1', <tensorflow.python.keras.wrappers.scikit_learn.KerasClassifier object at 0x7f4ed7d4c518>), ('model2', <tensorflow.python.keras.wrappers.scikit_learn.KerasClassifier object at 0x7f4ed7d4c470>), ('model3', <tensorflow.python.keras.wrappers.scikit_learn.KerasClassifier object at 0x7f4ed7d4c588>)],
    flatten_transform=None, n_jobs=None, voting='soft', weights=None)
```

```
In [57]: y_pred = ensemble_clf.predict(x_test)
print('acc: ', accuracy_score(y_pred, y_test))
```

acc: 0.9504

9.全部使用

```
In [5]: from tensorflow.keras import layers

import numpy as np
from tensorflow.keras.wrappers.scikit_learn import KerasClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.metrics import accuracy_score

def mlp_model():
    model = keras.Sequential([
        layers.Dense(64, activation='relu', kernel_initializer='he_normal', input_shape=(1, 1000)),
        layers.BatchNormalization(),
        layers.Dropout(0.2),
        layers.Dense(64, activation='relu', kernel_initializer='he_normal'),
        layers.BatchNormalization(),
        layers.Dropout(0.2),
        layers.Dense(64, activation='relu', kernel_initializer='he_normal'),
        layers.BatchNormalization(),
        layers.Dropout(0.2),
        layers.Dense(10, activation='softmax')
    ])
    model.compile(optimizer=keras.optimizers.SGD(),
                  loss=keras.losses.SparseCategoricalCrossentropy(),
                  metrics=['accuracy'])
    return model

model1 = KerasClassifier(build_fn=mlp_model, epochs=100, verbose=0)
model2 = KerasClassifier(build_fn=mlp_model, epochs=100, verbose=0)
model3 = KerasClassifier(build_fn=mlp_model, epochs=100, verbose=0)
model4 = KerasClassifier(build_fn=mlp_model, epochs=100, verbose=0)
ensemble_clf = VotingClassifier(estimators=[
    ('model1', model1), ('model2', model2), ('model3', model3), ('model4', model4),
```

```
In [9]: ensemble_clf.fit(x_train, y_train)
```

```
Out[9]: VotingClassifier(estimators=[('model1', <tensorflow.python.keras.wrappers.scikit_learn.KerasClassifier object at 0x7f7183d246a0>), ('model2', <tensorflow.python.keras.wrappers.scikit_learn.KerasClassifier object at 0x7f7183d245c0>), ('model3', <tensorflow.python.keras.wrappers.scikit_learn.KerasClassifier object at 0x7f7183d5a198>), ('model4', <tensorflow.python.keras.wrappers.scikit_learn.KerasClassifier object at 0x7f71839c94e0>)],
                        flatten_transform=None, n_jobs=None, voting='hard', weights=None)
```

```
In [10]: y_predict = ensemble_clf.predict(x_test)
```

```
In [11]: print('acc: ', accuracy_scoreaccuracy_scoreaccuracy_score(y_pred, y_test))
```

```
-----  
NameError                                Traceback (most recent call last)  
<ipython-input-11-c1dea66b1d3d> in <module>  
----> 1 print('acc: ', accuracy_score(y_pred, y_test))  
  
NameError: name 'y_pred' is not defined
```

```
In [ ]:
```