

tensorflow2教程-CNN變體網路

```
In [1]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
print(tf.__version__)
```

2.3.1

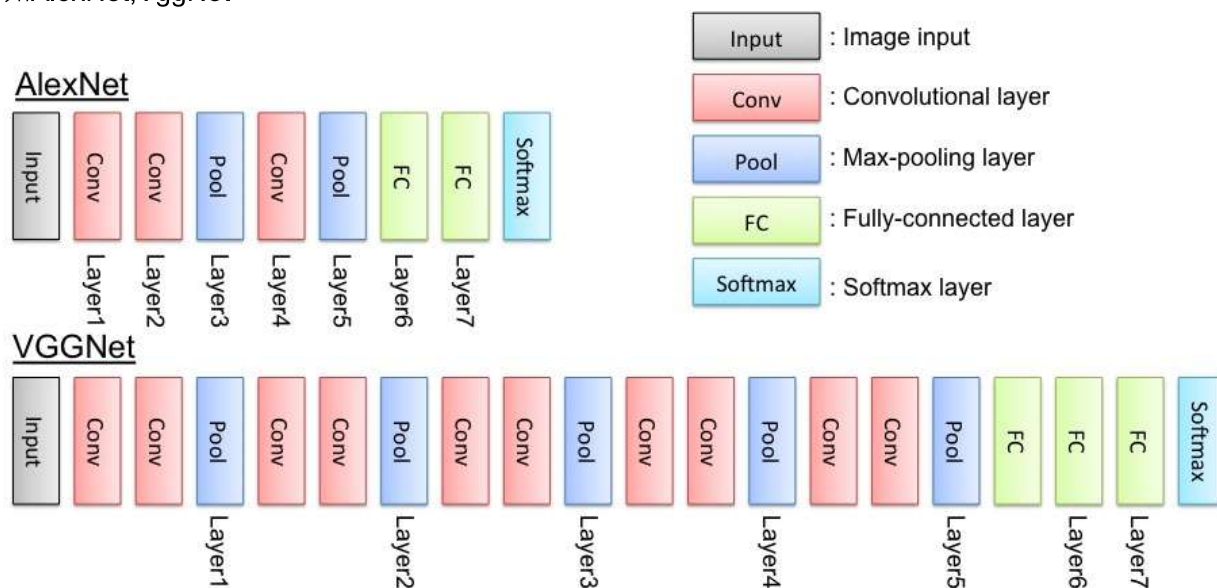
1.載入數據

```
In [2]: (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
x_train = x_train.reshape((-1,28,28,1))
x_test = x_test.reshape((-1,28,28,1))
print(x_train.shape, ' ', y_train.shape)
print(x_test.shape, ' ', y_test.shape)
```

```
(60000, 28, 28, 1) (60000,)
(10000, 28, 28, 1) (10000,)
```

2.簡單的深度網路

如AlexNet,VggNet



```
In [3]: x_shape = x_train.shape
deep_model = keras.Sequential(
[
    layers.Conv2D(input_shape=((x_shape[1], x_shape[2], x_shape[3])),
        filters=32, kernel_size=(3,3), strides=(1,1), padding='same', a
    layers.Conv2D(filters=32, kernel_size=(3,3), strides=(1,1), padding='same', a
    layers.MaxPool2D(pool_size=(2,2)),
    layers.Conv2D(filters=32, kernel_size=(3,3), strides=(1,1), padding='same', a
    layers.Conv2D(filters=32, kernel_size=(3,3), strides=(1,1), padding='same', a
    layers.MaxPool2D(pool_size=(2,2)),
    layers.Flatten(),
    layers.Dense(32, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

```
In [4]: deep_model.compile(optimizer=keras.optimizers.Adam(),
    loss=keras.losses.SparseCategoricalCrossentropy(),
    metrics=['accuracy'])
deep_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 32)	320
conv2d_1 (Conv2D)	(None, 28, 28, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 14, 14, 32)	9248
conv2d_3 (Conv2D)	(None, 14, 14, 32)	9248
max_pooling2d_1 (MaxPooling2	(None, 7, 7, 32)	0
flatten (Flatten)	(None, 1568)	0
dense (Dense)	(None, 32)	50208
dense_1 (Dense)	(None, 10)	330
Total params: 78,602		
Trainable params: 78,602		
Non-trainable params: 0		

```
In [5]: history = deep_model.fit(x_train, y_train, batch_size=64, epochs=5, validation_s
```


```
Epoch 1/5
844/844 [=====] - 65s 76ms/step - loss: 0.2987 - accur
acy: 0.9249 - val_loss: 0.1042 - val_accuracy: 0.9720
Epoch 2/5
844/844 [=====] - 70s 83ms/step - loss: 0.0657 - accur
acy: 0.9797 - val_loss: 0.0457 - val_accuracy: 0.9862
Epoch 3/5
844/844 [=====] - 75s 88ms/step - loss: 0.0463 - accur
acy: 0.9855 - val_loss: 0.0388 - val_accuracy: 0.9885
Epoch 4/5
844/844 [=====] - 86s 102ms/step - loss: 0.0371 - accu
racy: 0.9884 - val_loss: 0.0564 - val_accuracy: 0.9828
Epoch 5/5
844/844 [=====] - 92s 109ms/step - loss: 0.0335 - accu
racy: 0.9892 - val_loss: 0.0381 - val_accuracy: 0.9882
```

```
In [6]: deep_model.evaluate(x_test, y_test)
```

```
313/313 [=====] - 4s 14ms/step - loss: 0.0416 - accura
cy: 0.9882
```

```
Out[6]: [0.04156742990016937, 0.9882000088691711]
```

```
In [7]: import matplotlib.pyplot as plt
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'valivation'], loc='upper left')
plt.show()
```



```
In [10]: result = deep_model.evaluate(x_test, y_test)
```

```
10000/10000 [=====] - 2s 219us/sample - loss: 0.0445 -
accuracy: 0.9863
```

3. 添加了其它功能层的深度卷积

```
In [8]: x_shape = x_train.shape
deep_model = keras.Sequential(
[
    layers.Conv2D(input_shape=((x_shape[1], x_shape[2], x_shape[3])),
        filters=32, kernel_size=(3,3), strides=(1,1), padding='same', activation='relu'),
    layers.BatchNormalization(),
    layers.Conv2D(filters=32, kernel_size=(3,3), strides=(1,1), padding='same', activation='relu'),
    layers.BatchNormalization(),
    layers.MaxPool2D(pool_size=(2,2)),
    layers.Conv2D(filters=32, kernel_size=(3,3), strides=(1,1), padding='same', activation='relu'),
    layers.BatchNormalization(),
    layers.BatchNormalization(),
    layers.Conv2D(filters=32, kernel_size=(3,3), strides=(1,1), padding='same', activation='relu'),
    layers.MaxPool2D(pool_size=(2,2)),
    layers.Flatten(),
    layers.Dense(32, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(10, activation='softmax')
])
```

```
In [10]: deep_model.compile(optimizer=keras.optimizers.Adam(),
                             loss=keras.losses.SparseCategoricalCrossentropy(),
                             metrics=['accuracy'])
deep_model.summary()
```

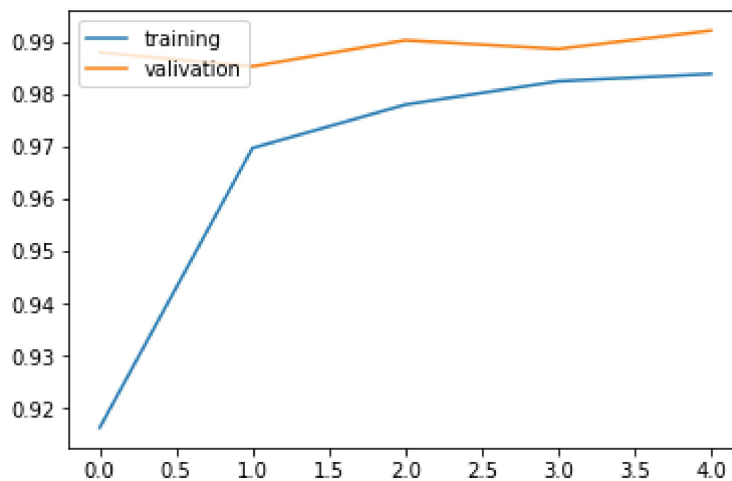
Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 28, 28, 32)	320
batch_normalization (BatchNo	(None, 28, 28, 32)	128
conv2d_5 (Conv2D)	(None, 28, 28, 32)	9248
batch_normalization_1 (Batch	(None, 28, 28, 32)	128
max_pooling2d_2 (MaxPooling2	(None, 14, 14, 32)	0
conv2d_6 (Conv2D)	(None, 14, 14, 32)	9248
batch_normalization_2 (Batch	(None, 14, 14, 32)	128
batch_normalization_3 (Batch	(None, 14, 14, 32)	128
conv2d_7 (Conv2D)	(None, 14, 14, 32)	9248
max_pooling2d_3 (MaxPooling2	(None, 7, 7, 32)	0
flatten_1 (Flatten)	(None, 1568)	0
dense_2 (Dense)	(None, 32)	50208
dropout (Dropout)	(None, 32)	0
dense_3 (Dense)	(None, 10)	330
=====		
Total params: 79,114		
Trainable params: 78,858		
Non-trainable params: 256		

```
In [11]: history = deep_model.fit(x_train, y_train, batch_size=64, epochs=5, validation_s
```

```
Epoch 1/5
844/844 [=====] - 161s 191ms/step - loss: 0.3188 - acc
uracy: 0.8983 - val_loss: 0.0598 - val_accuracy: 0.9822
Epoch 2/5
844/844 [=====] - 165s 195ms/step - loss: 0.1280 - acc
uracy: 0.9599 - val_loss: 0.0625 - val_accuracy: 0.9837
Epoch 3/5
844/844 [=====] - 182s 215ms/step - loss: 0.0981 - acc
uracy: 0.9693 - val_loss: 0.0368 - val_accuracy: 0.9897
Epoch 4/5
844/844 [=====] - 190s 225ms/step - loss: 0.0803 - acc
uracy: 0.9738 - val_loss: 0.0424 - val_accuracy: 0.9893
Epoch 5/5
844/844 [=====] - 181s 215ms/step - loss: 0.0662 - acc
uracy: 0.9789 - val_loss: 0.0412 - val_accuracy: 0.9890
```

```
In [14]: import matplotlib.pyplot as plt
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'valivation'], loc='upper left')
plt.show()
```




```
In [15]: result = deep_model.evaluate(x_test, y_test)
```

```
10000/10000 [=====] - 4s 365us/sample - loss: 0.0288 -
accuracy: 0.9909
```

4.NIN網路

Min等人在 2013年 (<https://arxiv.org/abs/1312.4400>) 提出了減少模型中參數數量的方法之一 (<https://arxiv.org/abs/1312.4400%EF%BC%89%E6%8F%90%E5%87%BA%E4%BA%86%E6%B8%>) 即“網路中的網路 (NIN) ”或“1X1卷積”方法很簡單 - 在其他卷積層之後添加卷積層 具有降低圖像空間的維度 (深度) 的效果，有效地減少了參數的數量



GoogleNet 中就用到了NIN結構

```
In [16]: x_shape = x_train.shape
deep_model = keras.Sequential(
[
    layers.Conv2D(input_shape=((x_shape[1], x_shape[2], x_shape[3])),
                  filters=32, kernel_size=(3,3), strides=(1,1), padding='same', activation='relu'),
    layers.BatchNormalization(),
    layers.Conv2D(filters=16, kernel_size=(1,1), strides=(1,1), padding='valid', activation='relu'),
    layers.BatchNormalization(),
    layers.MaxPool2D(pool_size=(2,2)),
    layers.Conv2D(filters=32, kernel_size=(3,3), strides=(1,1), padding='same', activation='relu'),
    layers.BatchNormalization(),
    layers.Conv2D(filters=16, kernel_size=(1,1), strides=(1,1), padding='valid', activation='relu'),
    layers.BatchNormalization(),
    layers.MaxPool2D(pool_size=(2,2)),
    layers.Flatten(),
    layers.Dense(32, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(10, activation='softmax')
])
```

```
In [17]: deep_model.compile(optimizer=keras.optimizers.Adam(),
                             loss=keras.losses.SparseCategoricalCrossentropy(),
                             metrics=['accuracy'])
deep_model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 28, 28, 32)	320
batch_normalization_v2_4 (Batch Normalization)	(None, 28, 28, 32)	128
conv2d_9 (Conv2D)	(None, 28, 28, 16)	528
batch_normalization_v2_5 (Batch Normalization)	(None, 28, 28, 16)	64
max_pooling2d_4 (MaxPooling2D)	(None, 14, 14, 16)	0
conv2d_10 (Conv2D)	(None, 14, 14, 32)	4640
batch_normalization_v2_6 (Batch Normalization)	(None, 14, 14, 32)	128
conv2d_11 (Conv2D)	(None, 14, 14, 16)	528
batch_normalization_v2_7 (Batch Normalization)	(None, 14, 14, 16)	64
max_pooling2d_5 (MaxPooling2D)	(None, 7, 7, 16)	0
flatten_2 (Flatten)	(None, 784)	0
dense_4 (Dense)	(None, 32)	25120
dropout_1 (Dropout)	(None, 32)	0
dense_5 (Dense)	(None, 10)	330
=====		
Total params: 31,850		
Trainable params: 31,658		
Non-trainable params: 192		


```
In [18]: history = deep_model.fit(x_train, y_train, batch_size=64, epochs=5, validation_sp
```

Train on 54000 samples, validate on 6000 samples

Epoch 1/5

54000/54000 [=====] - 62s 1ms/sample - loss: 0.2729 - accuracy: 0.9147 - val_loss: 0.0657 - val_accuracy: 0.9818

Epoch 2/5

54000/54000 [=====] - 63s 1ms/sample - loss: 0.0872 - accuracy: 0.9739 - val_loss: 0.0437 - val_accuracy: 0.9865

Epoch 3/5

54000/54000 [=====] - 59s 1ms/sample - loss: 0.0657 - accuracy: 0.9800 - val_loss: 0.0404 - val_accuracy: 0.9890


Epoch 4/5

54000/54000 [=====] - 49s 913us/sample - loss: 0.0535 - accuracy: 0.9834 - val_loss: 0.0622 - val_accuracy: 0.9830

Epoch 5/5

54000/54000 [=====] - 49s 913us/sample - loss: 0.0441 - accuracy: 0.9860 - val_loss: 0.0435 - val_accuracy: 0.9892

```
In [19]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'valivation'], loc='upper left')
plt.show()
```



```
In [20]: result = deep_model.evaluate(x_test, y_test)
```

10000/10000 [=====] - 2s 196us/sample - loss: 0.0335 - accuracy: 0.9887